

# CS156 (Introduction to AI), Spring 2022

## Homework 2 submission

Roster Name: Austin Rivard

Preferred Name (if different):

Student ID: 015044445

Email address: austin.rivard@sjsu.edu

Any special notes or anything you would like to communicate to me about this homework submission goes in here.

## References and sources

List all your references and sources here. This includes all sites/discussion boards/blogs/posts/etc. where you grabbed some code examples.

## Solution

Load libraries and set random number generator seed

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from scipy.spatial import distance
from sklearn.metrics import accuracy_score
```

```
In [2]: np.random.seed(42)
```

Code the solution

### 1) 2-D data

```
In [3]: def knn(newObservation, referenceData: pd.DataFrame, k=3):
# 1. create array of distances from each reference data point to new observati
# scipy.spatial.distance.euclidean()

distances = distance.cdist([newObservation], referenceData.values[:, 0:-1], me

# 2. nearestNeighbors are the points with least smallest distance, found using
```

```
nearestNeighbors = referenceData.assign(dist=distances[0]).nsmallest(k, 'dist')

# 3. the new observation is classified as the most common classification of it
# if there is a tie, take the first (but perhaps random would be better?)

return nearestNeighbors['classification'].mode().iloc[0]
```

In [4]:

```
#generate 2d data
n = 100

X1 = np.random.normal(loc=-2, scale=2, size=n//2)
X2 = np.random.normal(loc=2, scale=2, size=n//2)
X = np.concatenate((X1, X2))
Y = np.random.normal(size=n)
```

In [5]:

```
num_labels = 2
labels = np.concatenate([[i] * (n // num_labels) for i in range(num_labels)])
print(labels)
```

```
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0  
 0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1  
 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1]
```

In [6]:

```
df = pd.DataFrame({'X': X, 'Y': Y})
df.head()
```

Out[6]:

	X	Y
0	-1.006572	-1.415371
1	-2.276529	-0.420645
2	-0.704623	-0.342715
3	1.046060	-0.802277
4	-2.468307	-0.161286

In [7]:

```
X_train, X_test, Y_train, Y_test = train_test_split(df, labels, test_size=0.2, ra
```

In [8]:

```
ref_data = X_train.assign(classification = Y_train)
print(ref_data.head())
```

	X	Y	classification
43	-2.602207	0.184634	0
62	-0.212670	1.158596	1
3	1.046060	-0.802277	0
71	5.076073	-0.815810	1
45	-3.439688	0.781823	0

In [9]:

```
pred = knn(X_test.iloc[0], ref_data, k=3)
print(f'prediction = {pred}, actual = {Y_test[0]}')
```

```
prediction = 0, actual = 0
```

```
In [10]: Y_pred = X_test.apply(knn, axis=1, referenceData=ref_data, k=3)
print(Y_pred.head())
```

```
26    0
86    1
2     0
55    1
75    1
dtype: int64
```


```
In [11]: print(f'Accuracy of the predictions on the test data set is {accuracy_score(Y_te

plt.subplot(1, 2, 1)
plt.scatter(X_train.iloc[:,0],X_train.iloc[:,1], s=25, c=Y_train, marker=".")
plt.scatter(X_test.iloc[:,0],X_test.iloc[:,1], s=50, c=Y_test, marker="v")
plt.title("Actual labels")

plt.subplot(1, 2, 2)
plt.scatter(X_train.iloc[:,0],X_train.iloc[:,1], s=25, c=Y_train, marker=".")
plt.scatter(X_test.iloc[:,0],X_test.iloc[:,1], s=50, c=Y_pred, marker="v")
plt.title("Predicted labels")

plt.tight_layout()
plt.show()
```

Accuracy of the predictions on the test data set is 0.95



## 2) 3-D data

```
In [12]: #generate 3d data
n = 1000
num_classes = 4

X = np.random.normal(loc=0, scale=3, size=n)

Y1 = np.random.normal(loc=-3, scale=1, size=n//num_classes)
Y2 = np.random.normal(loc=1, scale=2, size=n//num_classes)
Y3 = np.random.normal(loc=3, scale=1, size=n//num_classes)
Y4 = np.random.normal(loc=5, scale=3, size=n//num_classes)
```

```
Y = np.concatenate((Y1, Y2, Y3, Y4))

Z1 = np.random.normal(loc=-1, scale=1, size=n//num_classes)
Z2 = np.random.normal(loc=1, scale=1, size=n//num_classes)
Z3 = np.random.normal(loc=4, scale=1, size=n//num_classes)
Z4 = np.random.normal(loc=-3, scale=1, size=n//num_classes)
Z = np.concatenate((Z1, Z2, Z3, Z4))

labels = np.concatenate([[i] * (n // num_classes) for i in range(num_classes)])
```

In [13]:

```
df = pd.DataFrame({'X': X, 'Y': Y, 'Z': Z})
df.head()
```

Out[13]:

	X	Y	Z
0	1.073362	-2.874775	0.804348
1	1.682354	-3.429406	-1.190904
2	3.249154	-2.877702	-0.280242
3	3.161406	-2.456702	-2.293273
4	-4.133008	-2.951140	-1.956436

In [14]:

```
X_train, X_test, Y_train, Y_test = train_test_split(df, labels, test_size=0.2, ra
```

In [15]:

```
ref_data = X_train.assign(classification = Y_train)
ref_data.head()
```

Out[15]:

	X	Y	Z	classification
687	-1.077876	3.424061	5.388338	2
500	-1.568169	3.350630	4.386809	2
332	-0.198239	1.646335	1.555513	1
979	-3.310768	5.020400	-2.484372	3
817	-1.451658	-1.011587	-5.832156	3

In [16]:

```
Y_pred = X_test.apply(knn, axis=1, referenceData=ref_data, k=3)
print(Y_pred)
```

```
993    3
859    3
298    1
553    2
672    1
..
679    2
722    2
215    0
653    2
```

```
150      0
Length: 200, dtype: int64
```

In [17]:

```
print(f'Accuracy of the predictions on the test data set is {accuracy_score(Y_te

fig = plt.figure(figsize=(18, 15))

ax = plt.subplot(2, 3, 1, projection='3d')
ax.scatter(X_train.iloc[:,0],X_train.iloc[:,1], X_train.iloc[:,2], s=25, c=Y_tra
ax.scatter(X_test.iloc[:,0],X_test.iloc[:,1], X_test.iloc[:,2], s=50, c=Y_test,
plt.title("Actual labels")

ax = plt.subplot(2, 3, 2, projection='3d')
ax.scatter(X_train.iloc[:,0],X_train.iloc[:,1], X_train.iloc[:,2], s=25, c=Y_tra
ax.scatter(X_test.iloc[:,0],X_test.iloc[:,1], X_test.iloc[:,2], s=50, c=Y_pred,
plt.title("Predicted labels")

plt.tight_layout()
plt.show()
```

Accuracy of the predictions on the test data set is 0.915

Actual labels

Predicted labels

